



Core C++ 2025

19 Oct. 2025 :: Tel-Aviv

From GPU Bottlenecks to Smooth Chat: Cost-Efficient Architectures for LLM Inference

Eshcar Hillel



OpenAI and NVIDIA Announce Strategic Partnership to Deploy 10 Gigawatts of NVIDIA Systems

September 22, 2025



C++ Devs for LLM inference

https://www.ted.com/talks/eric_schmidt_the_ai_revolution_is_underhyped

Solving LLM Infrastructure Scalability Challenges

Serve **MORE** queries
with **LARGER** models--
FASTER, under the
SAME POWER budget

Infrastructure Challenges

GPU supply shortage

COST

Power limit

Better perf/watt

LLM solution addresses the main scalability issues

Interactive experience

High-quality responses

Larger smarter models


Time-To-First-Token SLA

Time-Per-Output-Token SLA

Better UX/\$

User Experience

System Efficiency vs User Experience Tradeoff




Setup:

- HW: 8* H100, PCIe Gen 5, Dell server
- Model: llama-3-70B, 8FP, GQA 8, TP 2
- Prompt: 100-3500, output: 64
- Different batch sizes

Agenda for Today's Talk

- ❑ Why prefill and KV-cache dominate E2E performance
- ❑ KV-cache Offloading: Shared KV-store approach (Disaggregated)
- ❑ Results, theoretical model and implications for system design

Popular LLM Applications w Repetitive Context



Chatbots

reuse early chat
content as context
for later chat input

Bug Fixing


Frequently using
entire code repository
as context

Insights

Assistant and rec sys query
large document sets

**Repetitive computations * highly inefficient * up to 99% prefill cost *
limit HBM bandwidth utilization and e2e performance**


Prompt Computation vs. Token Generation



- **Prefill phase (prompt)**
 - All input tokens processed in parallel to generate the first output token
 - Time to first token (TTFT)
 - Compute bound
- **Decode phase (token)**
 - Serialized token generation
 - Time per output token (TPOT)
 - Memory bound


Multi-Turn Chats Evict KV-Cache Between Turns

Multi-turn conversation or multi-shot task agent prefill attention kv-cache based on expanding history




KV-Cache Offloading in a Nutshell

Multi-turn conversation or multi-shot task agent prefill attention kv-cache based on expanding history




IO-Compute Parallelism

Prefill attention KV-cache: (2) retrieve past (1) compute new interaction (3) store new




IO-Compute Parallelism – Nsight View

Layer wise pipelining: IOs execute in parallel with new prompt kv-cache computation



Stratified classification of KV-cache offloading (KV-CO) Solutions

- **S0:** Simple but tiny capacity
- **S1:** extended capacity, remains non-sharable, incurs higher cost
- **S2:** sharable namespace and deep at the expense of performance or cost
- **S3:** sharable namespace and deep with improved performance





Core C++ 2025

19 Oct. 2025 :: Tel-Aviv

Prefill Speedup Analysis

Prefill Acceleration By KV-Cache Retrieval

- R - time to retrieve a token from storage
 - $R = |KV_{token}| / BW_{IO}$
 - T - compute time per input token
 - $T = (2|model| + N_{in} \cdot O(\sqrt{|model|})) / TFLOPs(*)$
 - Effective token acceleration $e = T/R$
 - α – KV-cache fraction cached in storage
 - N_{in} – number of input token
- (*) approximate attention compute;
assuming all GPU compute is
utilized, not always true

Prefill Acceleration Analysis

- $TTFT^V = N_{in} \cdot T$
- Offloading writes the newly generated KV cache
 - Let W denote the write time of a single token KV
- Assuming IO and compute can be overlapped
 - At each layer prefetch next layer's KV cache and write the KV cache of the previous layer
 - $TTFT^{KV} = \max(\underbrace{\alpha \cdot N_{in} \cdot T / e}_{\text{IO time to fetch existing KV cache from storage}}, \underbrace{(1-\alpha) \cdot N_{in} \cdot W}_{\text{IO time to write the new part of the prompt}}, \underbrace{(1-\alpha) \cdot N_{in} \cdot T}_{\text{Compute time of the new part of the prompt}})$

IO time to fetch existing KV cache from storage IO time to write the new part of the prompt Compute time of the new part of the prompt

Prefill Acceleration Analysis (Cont.)

- $TTFT^{KV} = \max(\alpha \cdot N_{in} \cdot T / e, (1-\alpha) \cdot N_{in} \cdot W, (1-\alpha) \cdot N_{in} \cdot T)$
 - When $W \leq T$ TTFT is not affected by write IOs
- $TTFT^{KV} = \max(\alpha \cdot N_{in} \cdot T / e, (1-\alpha) \cdot N_{in} \cdot T)$
- $TTFT^V = N_{in} \cdot T$

Insight
Write only need to
match compute
performance

Performance gain (speedup)

- **compute-bound:** $x=1/(1-\alpha)$
- **IO-bound:** $x=e/\alpha$


Acceleration
depends on read
performance and
hit rate

Prefill Acceleration Analysis (Cont.)

For example $e=9$
 e is the acceleration
from offloading (T/R)

Crossover
point
 $\alpha^*=e/(1+e)$


Insight
No advantage for
higher IO speed!
(HBM, DRAM)




Prefill Acceleration Analysis (Cont.)

For example $e=9$
 e is the acceleration
from offloading (T/R)

Crossover
point
 $\alpha^*=e/(1+e)$




Prefill Acceleration - Compute-Bound Example



Replacing GPU compute with storage IO in prefill allows higher HBM BW efficiency in decode via larger batch size

- HW: H100 gen 5 Dell server
- Model: llama-3-70B, 8FP, GQA 8, TP 1
- SLA: TTFT 400ms, TPOT 40ms
- Prompt: 100-3500, output: 64
- Maximum batch size to meet SLA




Core C++ 2025

19 Oct. 2025 :: Tel-Aviv

End-to-End (E2E) Gain for Continuous Batching

Colocating Prefill and Decode




- In-flight batching (continuous)
- Decode “love” batches as it is memory bw bound
- Prefill-prefill interference
- Prefill-decode interference
- Creating “bubbles”



E2E Gain - Continuous Batching Analysis

- B – number of concurrent requests
- N_{out} – Average number of output tokens
- Number of concurrent prefills $\sim Binom(B, 1/N_{out})$




E2E Gain - Continuous Batching Analysis

- B – number of concurrent requests
- N_{out} – Average number of output tokens
- Number of concurrent prefills $\sim Binom(B, 1/N_{out})$
- $E[TTFT(B)] = TTFT(1) + (B-1)/N_{out} \cdot TTFT(1) = (1 + (B-1)/N_{out}) \cdot TTFT(1)$
- $E[TPOT(B)] = TPOT(1) + (B-1) \cdot (TTFT(1)/N_{out} + |KV(1)|/BW_{HBM})$
 - $|KV(1)|$ is the average KV cache size of a single input prompt
 - BW_{HBM} is the memory BW

Expected number of additional simultaneous prefills in a prefill slot

Time to transfer a single prompt KV cache to compute engines

Measured TPOT Overhead for LlaMa-3-70B



Impact of Prefill Speedup on Decode Speedup

- $TPOT(B) \approx TPOT(1) + (B-1) \cdot (\Delta P + \Delta H)$
- $TPOT(B) \leq SLA_{TPOT}$
- Maximal B to meet SLA_{TPOT} : $\lfloor (SLA_{TPOT} - TPOT(1)) / (\Delta P + \Delta H) \rfloor$
- $\Delta P_{KV} = \Delta P/x$
- $B_{KV}/B_{vanilla} = (\Delta P + \Delta H) / (\Delta P_{KV} + \Delta H)$
 $= (\Delta P + \Delta H) / (\Delta P/x + \Delta H)$
- When $x \rightarrow \infty$, $B_{KV}/B_{vanilla} = 1 + \Delta P/\Delta H$

System is required to meet SLA

x is the speedup of the prefill using IO

Insight
TPS gain

Asymptotic E2E Gain Analysis

- TPS gain $\leq 1 + \Delta P / \Delta H$ aim to maximize $\Delta P / \Delta H$
- $\Delta P = TTFT(1) / N_{out}$, $\Delta H = |KV(1)| / BW_{HBM}$
 - $TTFT(1) = N_{in} \cdot (\underbrace{2|model| / TFLOPs}_{\text{“Model” Compute Time}} + \underbrace{N_{in} \cdot O(\sqrt{|model|}) / TFLOPs}_{\text{Attention Compute Time}}))$
- $\Delta H = \underbrace{N_{in} |KV_{token}|}_{|KV(1)|} / BW_{HBM}$

KV cache size of
a single token

Asymptotic E2E Gain Analysis (Cont.)


TTFT(1) behaves differently in two distinct regimes



- Short prompt regime : $\Delta P / \Delta H \propto |model| / |KV_{token}| \cdot BW_{HBM} / TFLOPs \cdot 1 / N_{out}$
 $(N_{in} \ll 6d)$
- Long prompt regime : $\Delta P / \Delta H \propto \sqrt{|model| / |KV_{token}|} \cdot BW_{HBM} / TFLOPs \cdot N_{in} / N_{out}$
 $(N_{in} \gg 6d)$

Pliops/Local DRAM/Local SSD Efficiency Comparison

System Efficiency vs User Experience Tradeoff: Pliops vs LMCache



- HW: H100 gen 5 Dell server
- Model: Qwen3-14B, 16FP
- Prompt: 9600 (avg), 6.5 turns (avg)
- output: 80-120
- Different batch sizes




Core C++ 2025

19 Oct. 2025 :: Tel-Aviv

AI Factories
Data-Center Scale LLM Inference

Datacenter Scale LLM Inference Framework




- Dynamo (Nvidia)
- lIbm-d (IBM/RedHat/Google)
- AI Brix (ByteDance)
- Production Stack (Tensor Mesh)

Support for Prefill-Decode Disaggregation (PDD)

Simplicity for Performance and Robustness

Monitor and map all events of creation and eviction of kv-cache entries


KV-cache aware routing
Full KV-Cache mapping





- ✗ Storage capacity coupled w compute (local SSD)
- ✗ Hinders scalability (centralized management)
- ✗ Memory footprint (global kv-cache map)
- ✗ Complex routing policies

KV-cache oblivious routing


- ✓ Unlimited capacity
- ✓ Offloading kv pool index management, fine granularity
- ✓ Supports elasticity, crash recovery, PD reconfigurations
- ✓ Lighter resource-focused routing policy



KV-Cache Offloading- System View




System Efficiency vs User Experience Tradeoff



- HW: H100 gen 5 Dell server
- Model: llama-3-70B, 8FP, GQA 8, TP 2
- Prompt: 2200, output: 170, turns: 15
- Different number of workers, clients

Summary & Conclusions

- ✓ KV-cache offloading increases efficiency and reduces cost while maintaining user experience constraints
- ✓ Throughput gains depend on the model, the GPU, and the workload
- ✓ In compute-bound cases, no benefit for faster-than-SSD memory
- ✓ Speedup depends mainly on read performance
- ✓ GPU-initiated IO achieves full IO-compute overlap with zero CPU overhead



Core C++ 2025

19 Oct. 2025 :: Tel-Aviv

THANK YOU

QUESTION



eshcar@pliops.com

linkedin.com/in/eshcar

