



Core C++ 2025

19 Oct. 2025 :: Tel-Aviv

“Lock me up, Scotty!”

Ben Liderman & Igor Khanin



About Us



Ben Liderman

System Architect,
Security & Trust Products



Igor Khanin

Senior Core Engineer,
Wallet Solutions

About Fireblocks

Fireblocks powers companies of all sizes to confidently build, run, and grow their business on the blockchain.

\$10T+

Transfers Secured

100+

Blockchains

300M+

Wallets Created

2,000+

Global Customers



Revolut





Reminder: The Mutual Exclusion Problem

Given **N** concurrent threads and a critical section **C**:

- **Safety** Property - at most one thread inside C
- **Liveness** Property - one of the following:



No Deadlocks

No Starvation / Basic Fairness

Bounded Wait / Strong Fairness

(As per prof. Hagit Attiya)



C++ has a Built-in Solution for the MEP

```
static std::mutex transporter_lock;
static Engineer scotty;

void beam_up_crew(const std::string& crew_member)
{
    std::lock_guard lg { transporter_lock };

    scotty.beam_me_up(crew_member);
    std::cout << crew_member
                  << " back on the Enterprise!"
                  << std::endl;
}
```



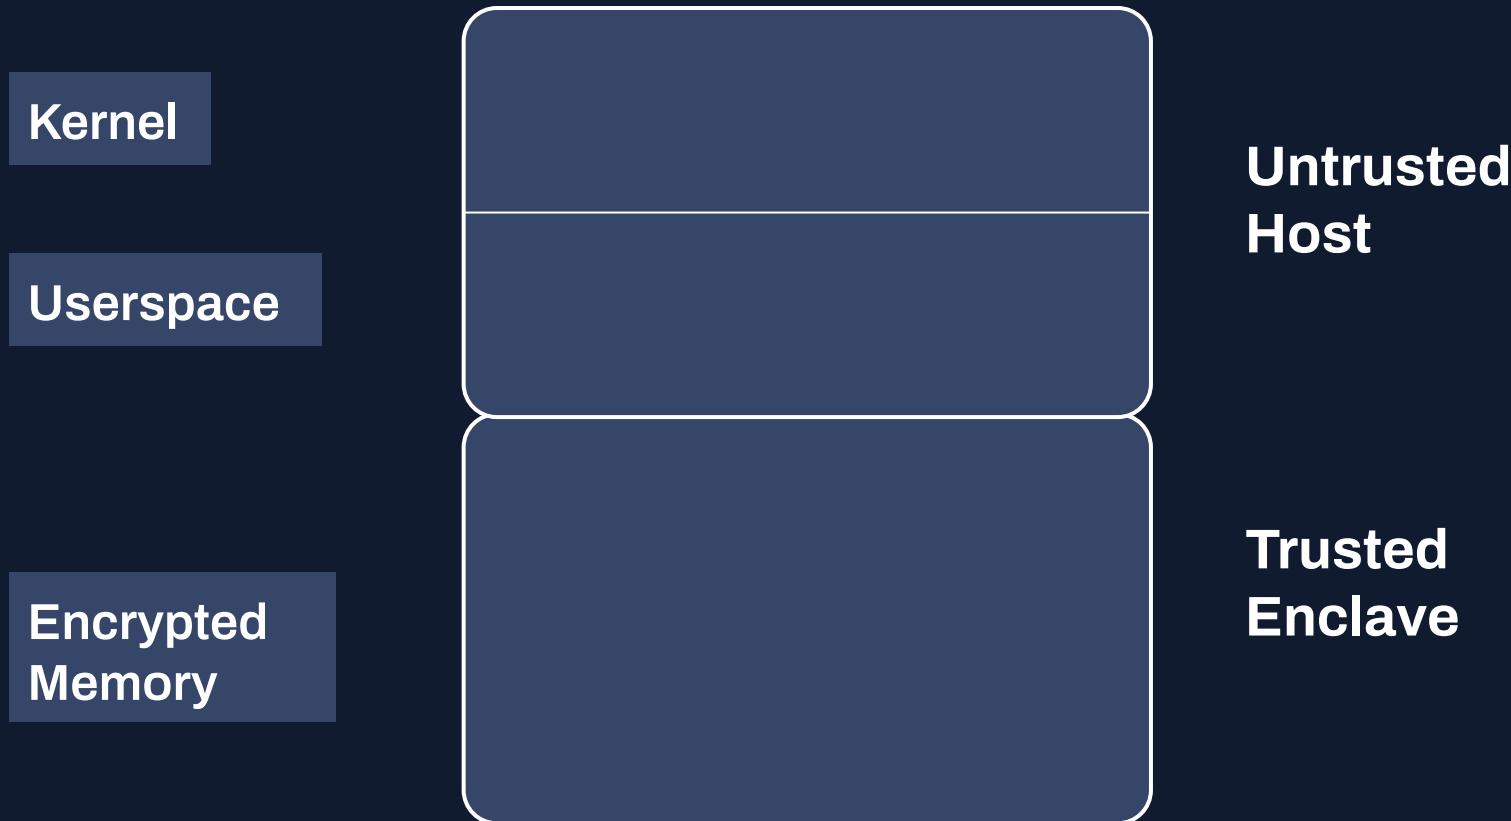
It Worked, Until it didn't...

CPU Usage by Container





Software Guard eXtensions from 10,000 feet



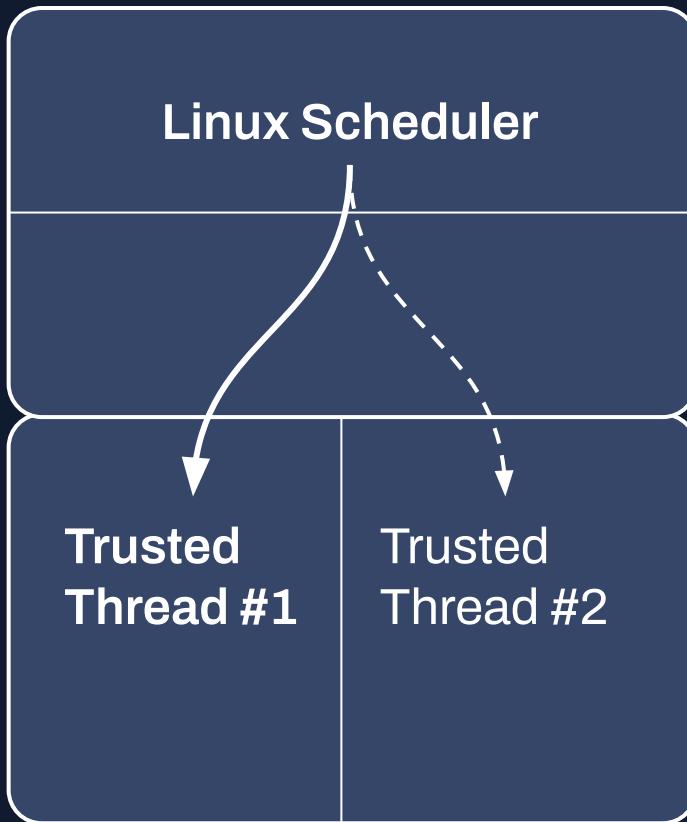


Software Guard eXtensions from 10,000 feet

Kernel

Userspace

Encrypted
Memory



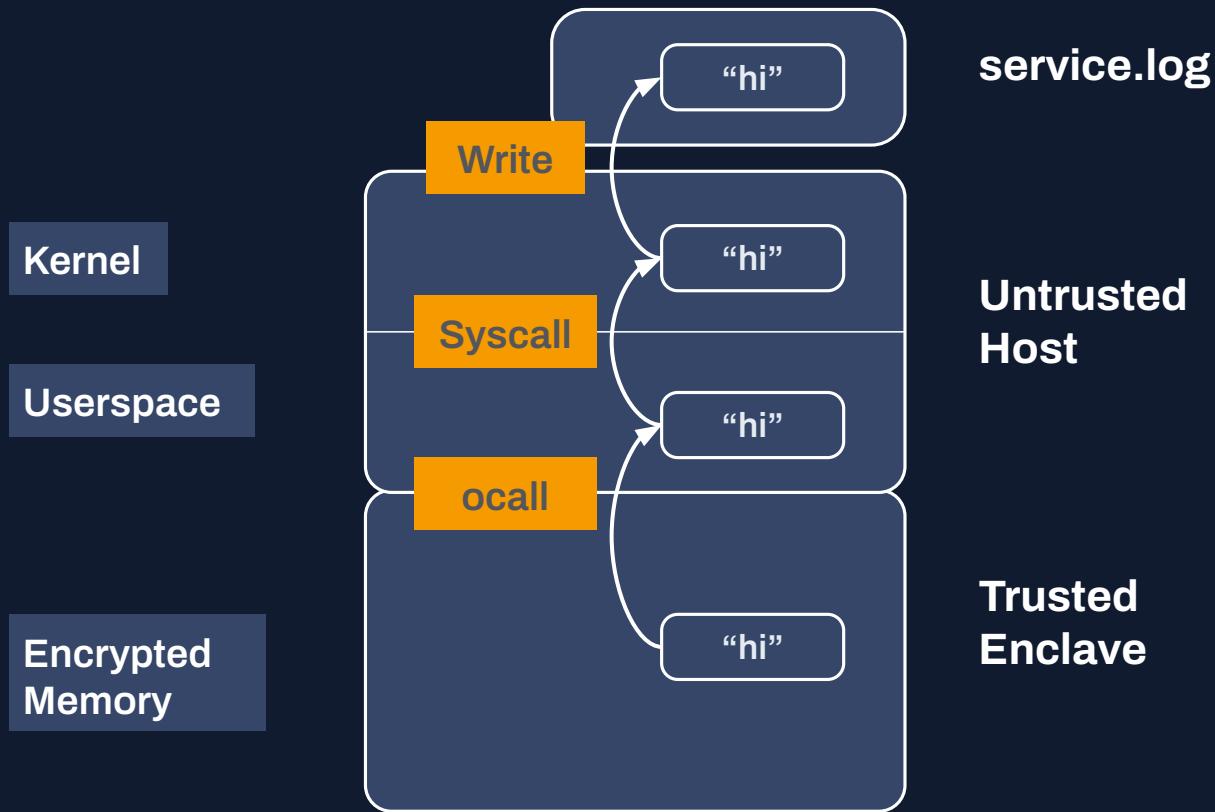
Untrusted
Host

Trusted
Enclave

ocall



Software Guard eXtensions from 10,000 feet





Sources of Mutual Exclusion for Applications

- Hardware support
- Operating system support
- Pure software algorithms (not practical)

▲ Mutual Exclusion with Hardware Support

```
std::atomic_flag is_locked = ATOMIC_FLAG_INIT;

void lock() {
    while (is_locked.test_and_set(std::memory_order_acquire))
    {
        // Spin!
    }
}

void unlock() {
    is_locked.clear(std::memory_order_release);
}
```



Mutual Exclusion with Hardware Support

- ✓ Safe and deadlock-free
- ✓ All state is inside the enclave
- ✓ Can be good under low contention
- ✗ Starvation
- ✗ Spinning wastes CPU
- ✗ Atomic access has overhead
- ✗ Breakdown under preemptive multitasking



Mutual Exclusion with OS Support (Linux)

```
constexpr uint32_t UNLOCKED = 0;
constexpr uint32_t LOCKED = 1;

alignas(4) uint32_t is_locked = UNLOCKED;
std::atomic_ref<uint32_t> is_locked_ref { is_locked };

void lock() {
    while (is_locked_ref.exchange(LOCKED, std::memory_order_acquire)
        != UNLOCKED) {
        syscall(SYS_futex, &is_locked, FUTEX_WAIT_PRIVATE, LOCKED, NULL);
    }
}

void unlock() {
    is_locked_ref.store(UNLOCKED, std::memory_order_release);
    syscall(SYS_futex, &is_locked, FUTEX_WAKE_PRIVATE, 1); // Wake one
}
```



A Balanced Approach

- We can mix the approaches
- Separate managing the lock state from thread control
- The “Parking Lot” design pattern



What we had

Critical Section
Control

```
int sgx_thread_mutex_lock(sgx_thread_mutex_t *mutex) {  
    while (1) {  
        SPIN_LOCK(&mutex->m_lock);  
        ...  
        if /*mutex is available*/ {  
            ...  
            mutex->m_owner = self;  
            SPIN_UNLOCK(&mutex->m_lock);  
            return 0;  
        }  
        QUEUE_INSERT_TAIL(&mutex->m_queue, self);  
        SPIN_UNLOCK(&mutex->m_lock);  
    }  
}
```

Thread Control

```
{  
    sgx_thread_wait_untrusted_event_ocall(self);  
}
```





A First Improvement

- We can eliminate the inner spinlock by placing the state in a single trusted atomic variable
- But it doesn't mean that we shouldn't spin at all!
- **Adaptive Locking:** Try to set the lock state N times before asking the operating system to suspend the thread



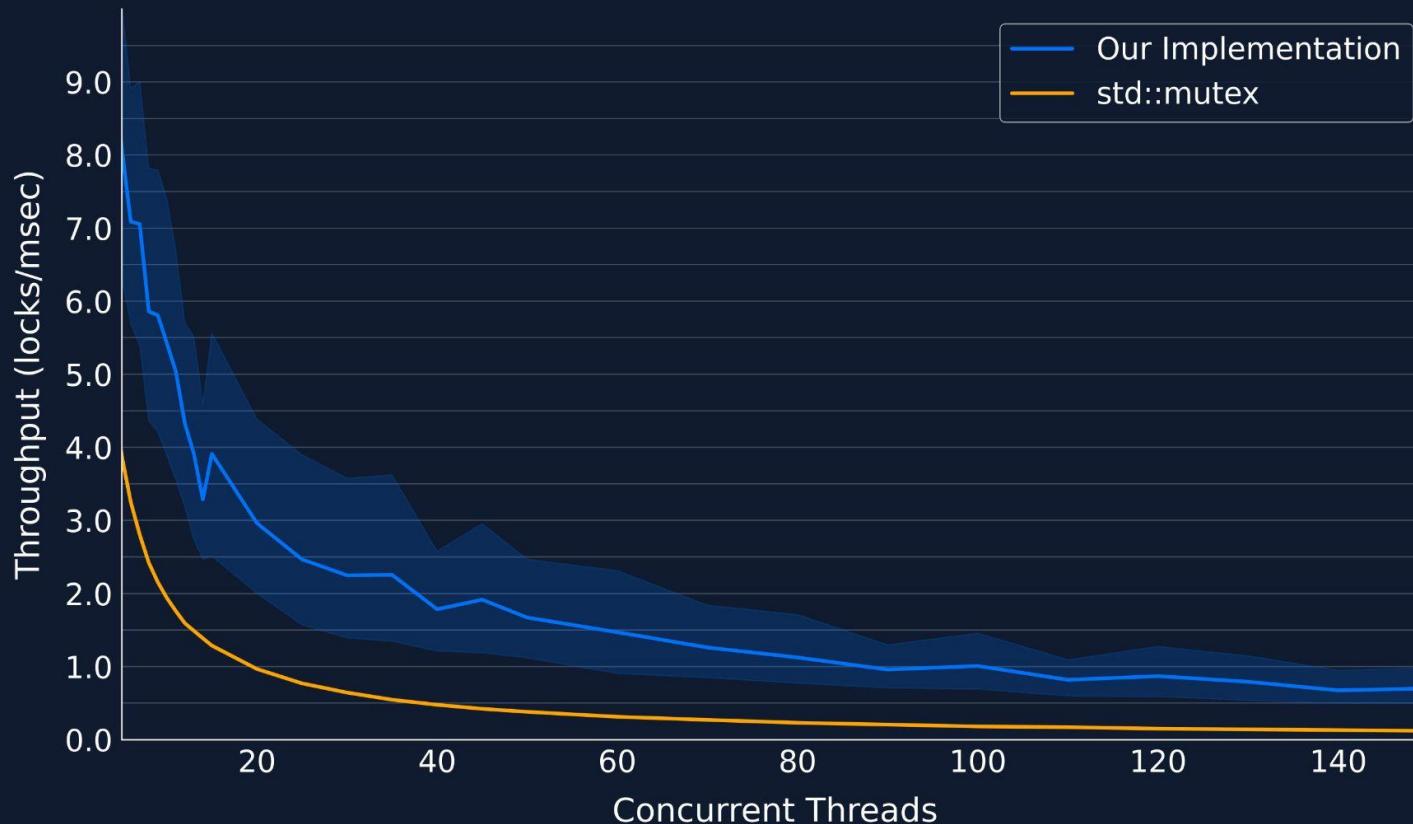
Is Fairness Good?

- Satisfying a strong liveness property is a trade-off
- Strong fairness reduces variance of time for entering the critical section, at the cost of both throughput and average latency
- Most platforms do not have a fairness guarantee by default



What we Got

Per-thread Throughput: Higher is Better





Should I do that too?

- “Locking isn’t slow, contention is”
- Primary effort should be directed at architecture choices that reduce contention
- Understand the guarantees that your specific platform is giving, and compare them to your workloads
- “Insufficient facts always invite danger”



Thank you!

