

Core C++ 2025

19 Oct. 2025 :: Tel-Aviv

(Intel)ligent Verification: Applying AI to Tame Microprocessor Complexity

Lotem Dalal

Lotem Dalal

- Curious
- Lifelong learning advocate
- Science and technology oriented

Agenda

- Introduction - CPU R&D cycle and verification
- why using AI for CPU verification
- Main uses of AI in CPU verification @ intel

Introduction

Designing a CPU core is like coordinating a million tiny, autonomous teams - each with different clocks, priorities and failure modes - to produce one perfectly ordered stream of instructions.
Proving it never fails, for decades.



introduction

- Micro-processor consists mainly of:

- **SoC** (system on chip) – peripheral subsystems
(GPU, I/O controllers, memory controller, accelerators)
- **CPU** (central processing unit) – processing
cores
(execution units, caches, control)
- **FW** (firmware) – hardware-software interface
embedded programmable SW layer (μ code, BIOS)

SoC

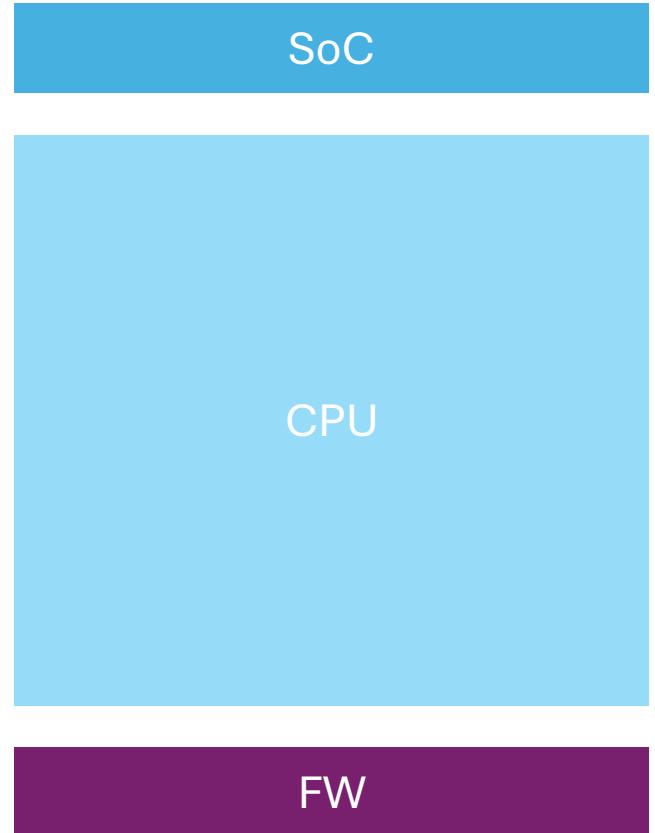
CPU

FW

CPU CORE R&D cycle

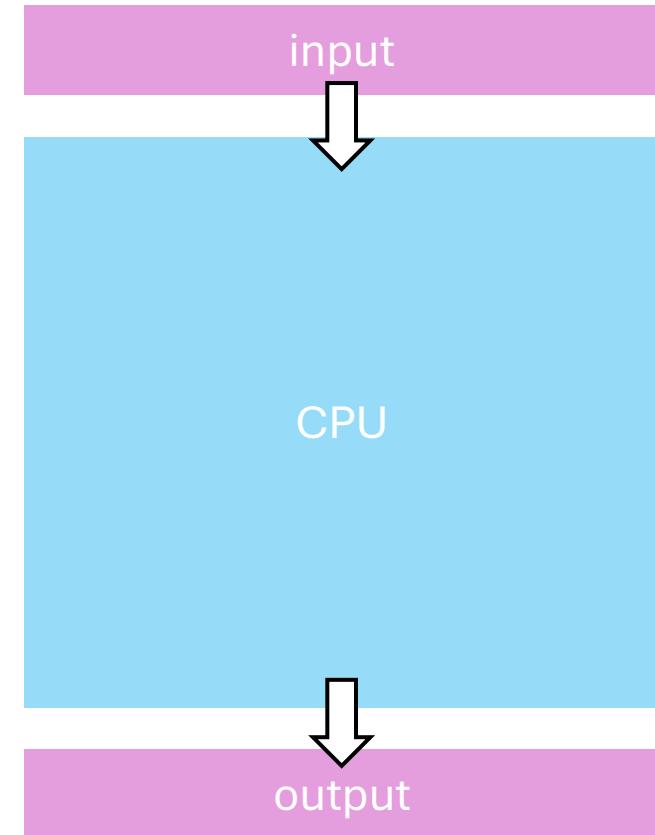
- Core of a CPU is the most complex IP, with:

- **2M+** Files
- **500M+** transistors
- **1k+** Engineers developing it



CPU CORE Verification

- Chip verification aims assuring no bugs are found in the DUT (**design under testing**).
- The methodology to achieve it consists of:
 - **Stimuli** – injecting constrained-random inputs.
 - **Checking** – rules the DUT must follow
 - **Coverage** - quantitative indicators that show how much of the design or specification has been verified.



CPU CORE Verification - overview

- Each project includes the following stages:

1. Planning 🔎

- Architectural spec review
- Design & Verification plans

2. Execution 💻

- Introducing new features
- Updating existing features

3. Analysis 🐞

- Focus on quality – debug and coverage

4. Closure ⚖️

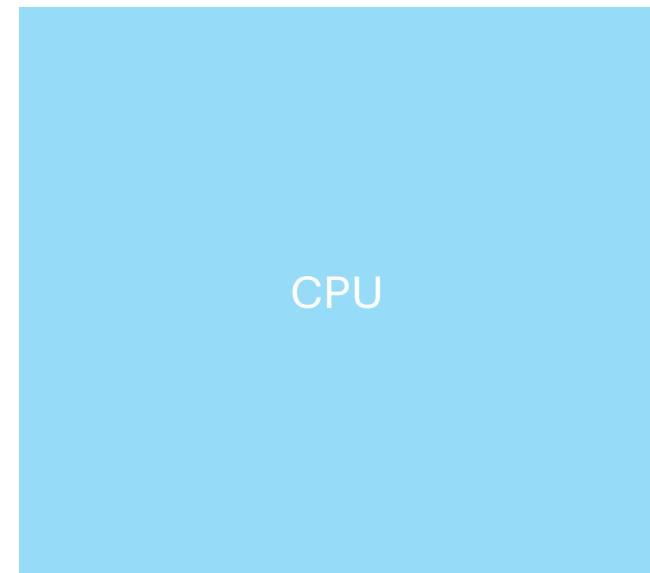
- Closing coverage gaps
- Bug fix is limited



CPU CORE Verification - overview

- Average time estimation of the stages:

1. **Planning** 🔎 10% (*dynamic)
 - Architectural spec review
 - Design & Verification plans
2. **Execution** 💻 70%
 - Introducing new features
 - Updating existing features
3. **Analysis** 🐞 15%
 - Focus on quality – debug and coverage
4. **Closure** ⚖️ 5%
 - Closing coverage gaps
 - Bug fix is limited



Most of the CPU (pre-Si) projects take ~2 years

CPU CORE Verification – domain level

- Each project includes the following phases:

1. Planning 🔎

- Architectural spec review
- Design & Verification plans

2. Execution 💻

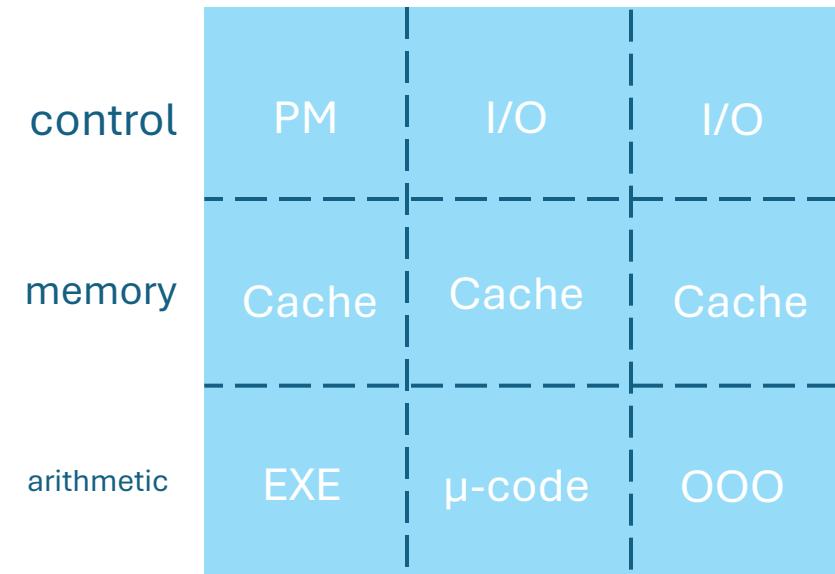
- Introducing new features
- Updating existing features

3. Analysis 🐞

- Focus on quality – debug and coverage

4. Closure ⚖️

- Closing coverage gaps
- Bug fix is limited



Most of the verification is done in a domain level

CPU CORE Verification – domain level

- Each project includes the following phases:

1. **Planning** 🔎

- Architectural spec review
- Design & Verification plans

2. **Execution** 💻

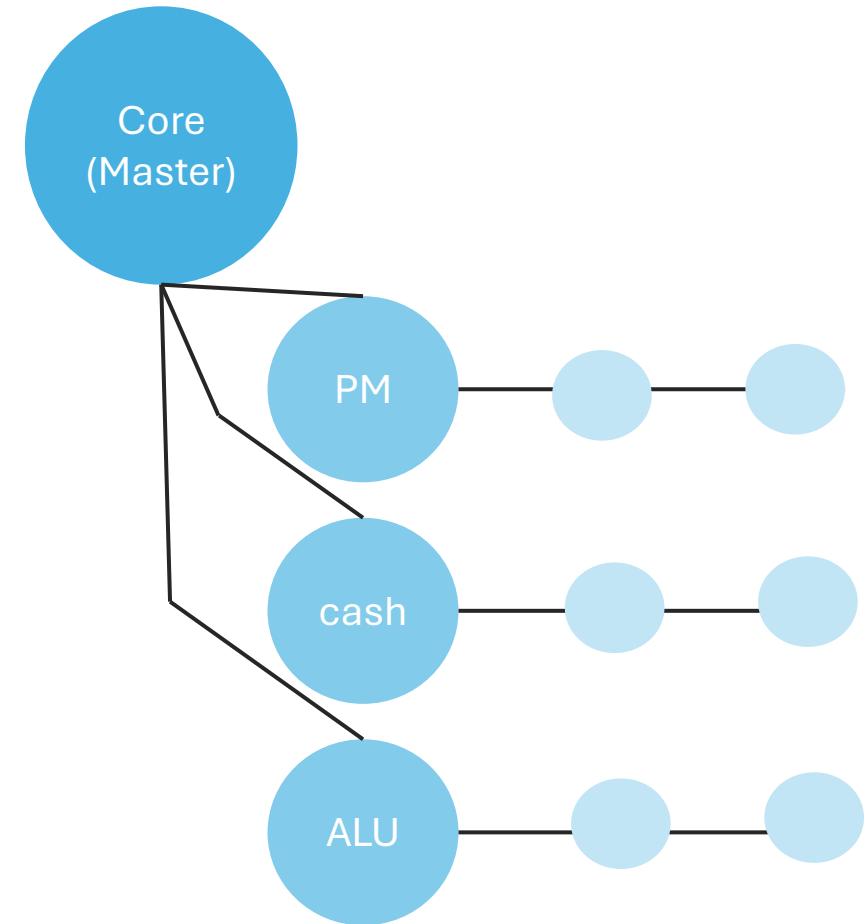
- Introducing new features
- Updating existing features

3. **Analysis** 🐞

- Focus on quality – debug and coverage

4. **Closure** ⚖️

- Closing coverage gaps
- Bug fix is limited



Most of the verification is done in a domain level

CPU CORE Verification – core level

- Each project includes the following phases:

1. Planning 🔎

- Architectural spec review
- Design & Verification plans

2. Execution 💻

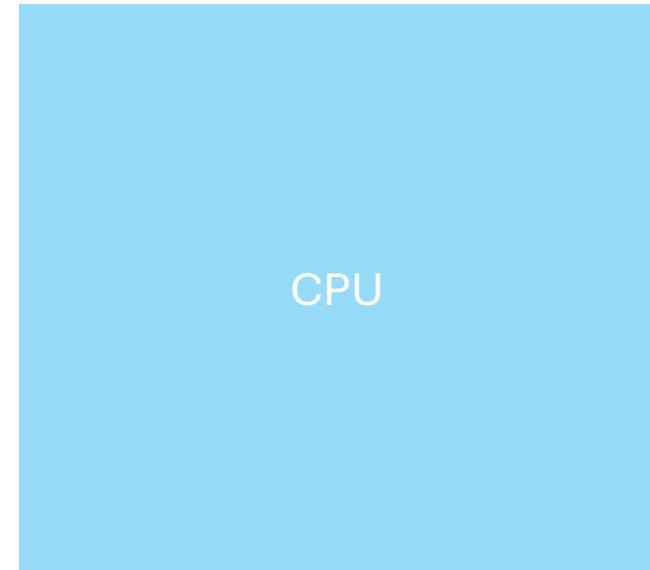
- Introducing new features
- Updating existing features

3. Analysis 🐞

- Focus on quality – debug and coverage

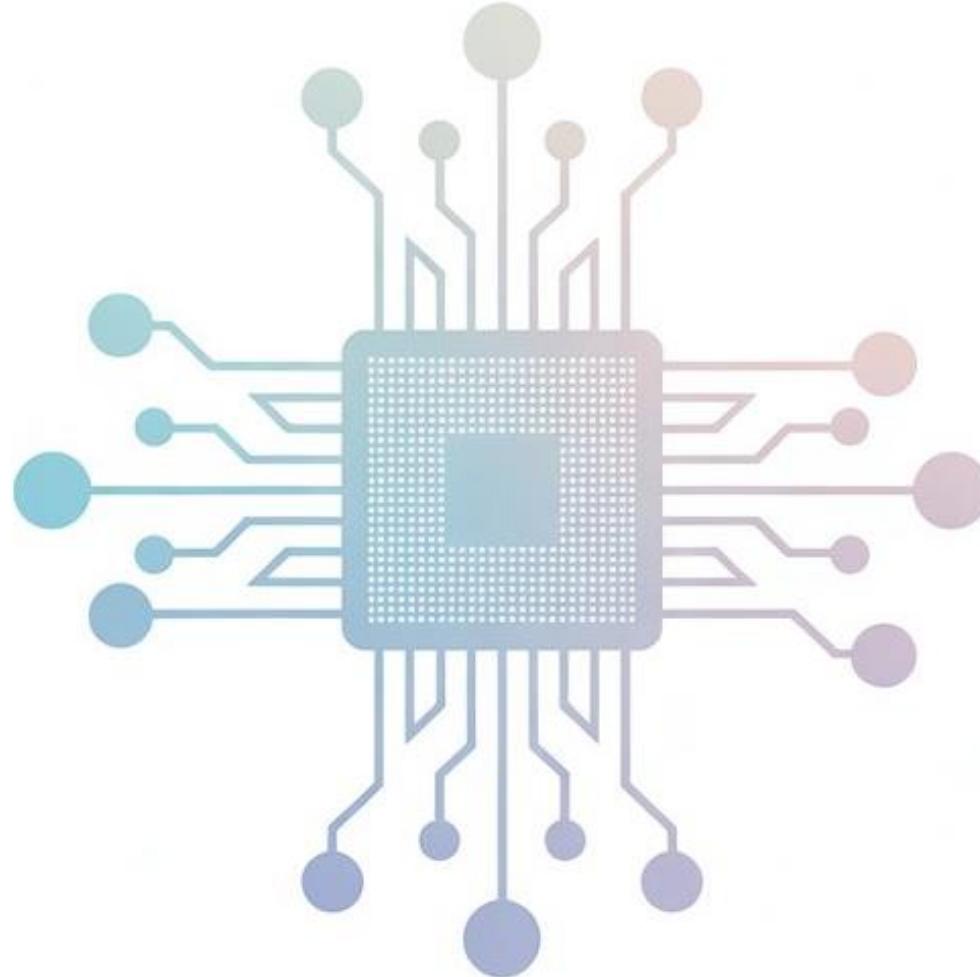
4. Closure ⚖️

- Closing coverage gaps
- Bug fix is limited



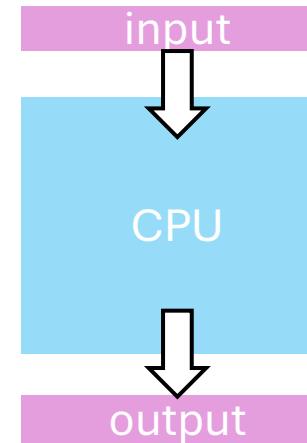
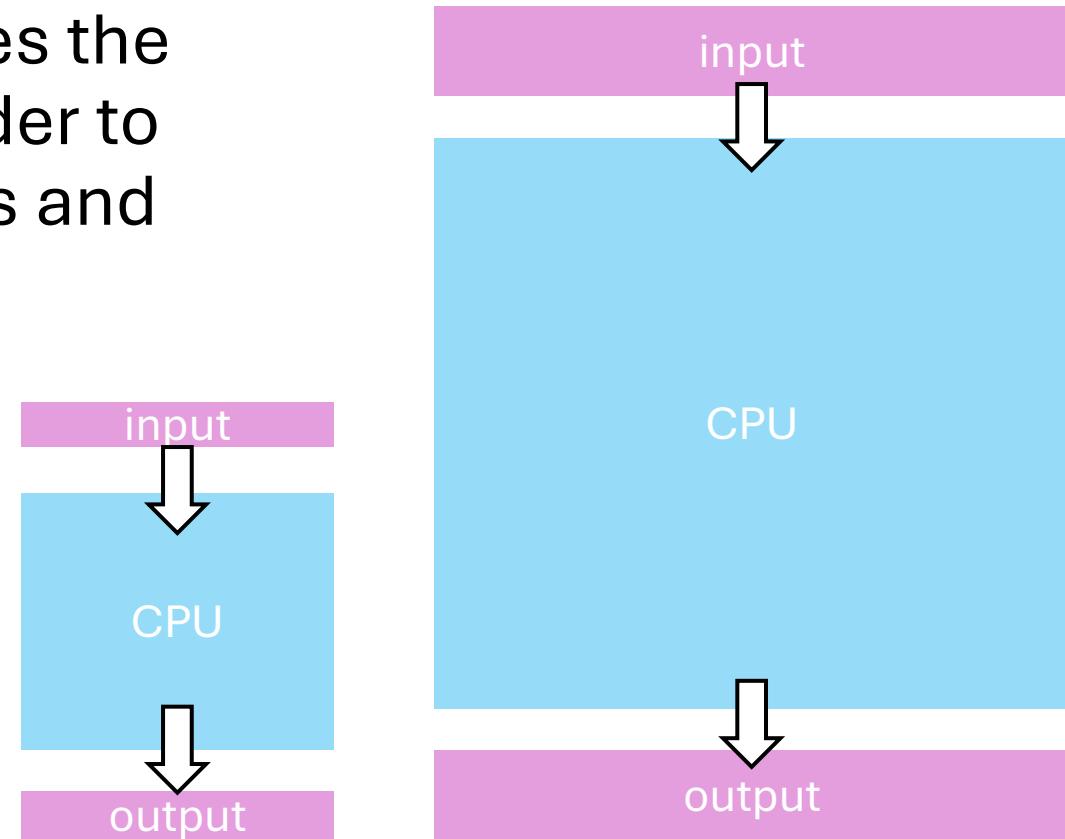
Main functionalities are also done in a **core level**

How can AI enhance CPU R&D?



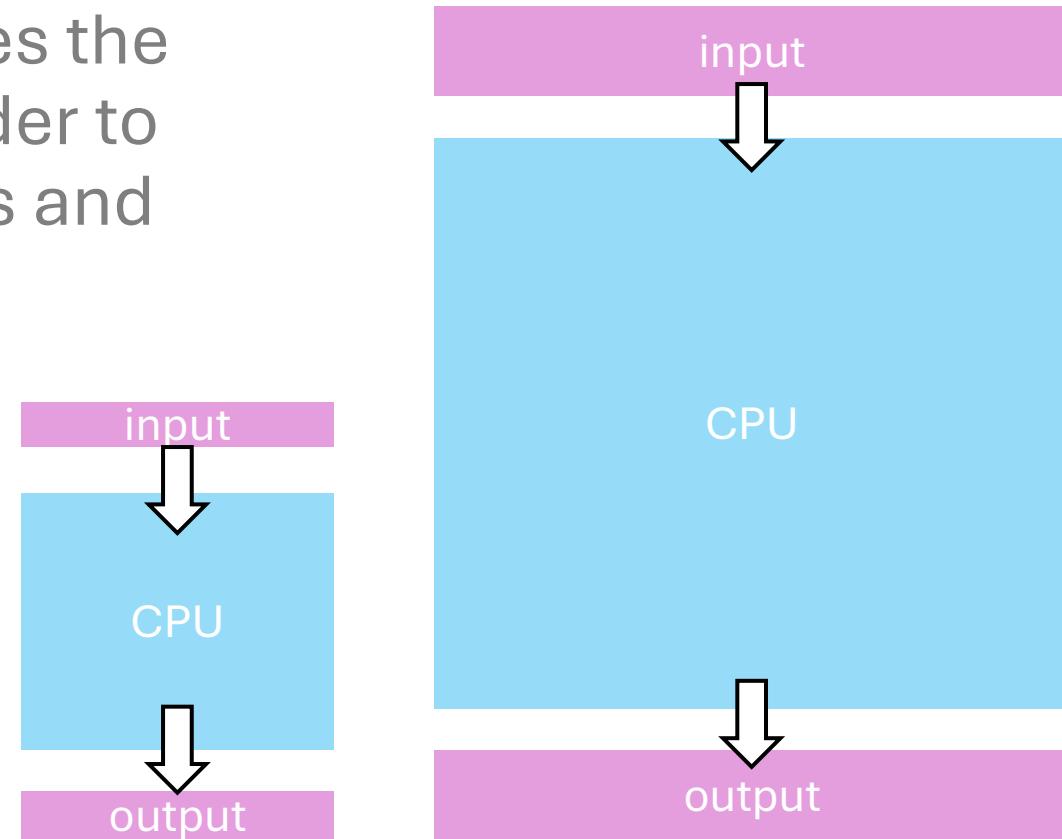
How can AI enhance CPU R&D?

- As technology advances, so does the CPU complexity – making it harder to simulate ALL possible scenarios and getting FULL coverage.



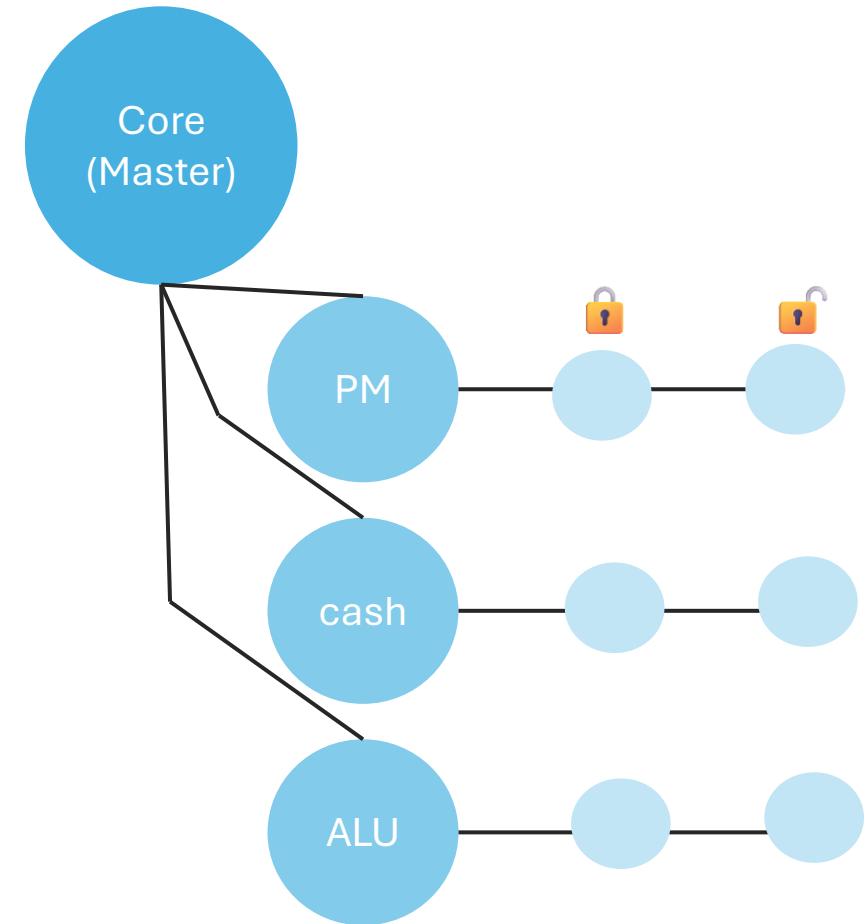
How can AI enhance CPU R&D?

- As technology advances, so does the CPU complexity – making it harder to simulate ALL possible scenarios and getting FULL coverage.
- Using AI and ML in CPU verification allows us to make the best out of the R&D resources.



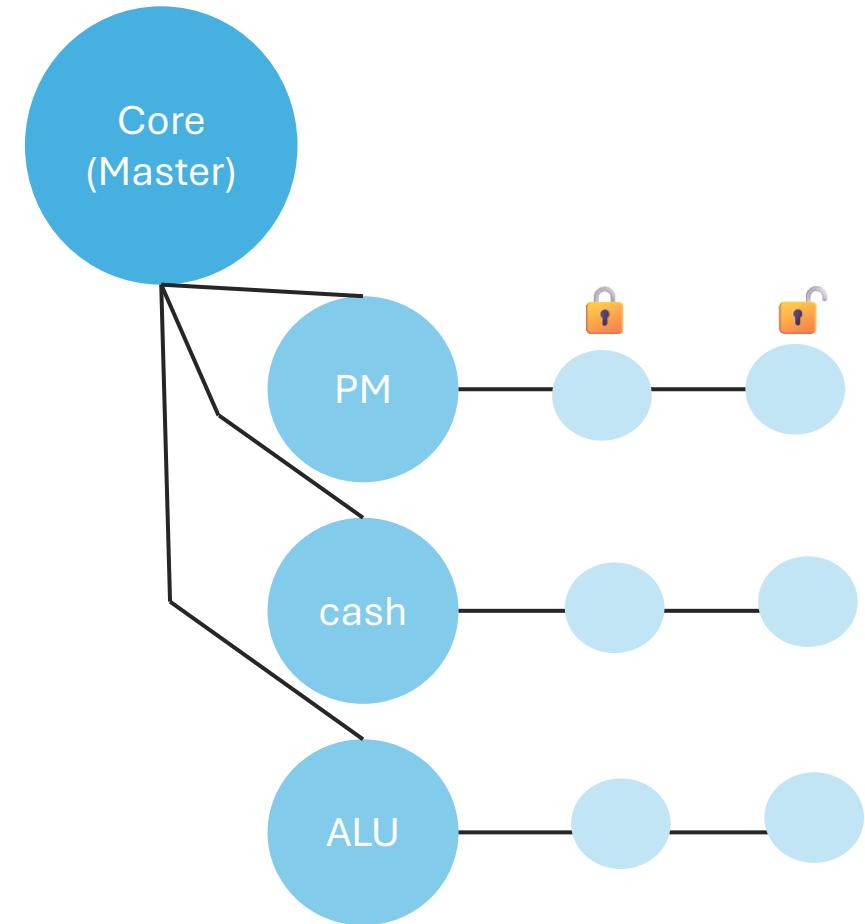
1. Optimizing gating lists

- “gate keeper” – identifying level 0 errors/bugs before merging to repo
 - Cover wide range of scenarios
 - **Not randomized**
 - Save compute resources
 - Minimum runtime – enhance R&D productivity



1. Optimizing gating lists

- “gate keeper” – identifying level 0 errors/bugs before merging to repo
- How can AI help me?
 - Analyze already written tests
 - Suggest minimum tests that cover maximum scenarios
 - Improving runtime of tests

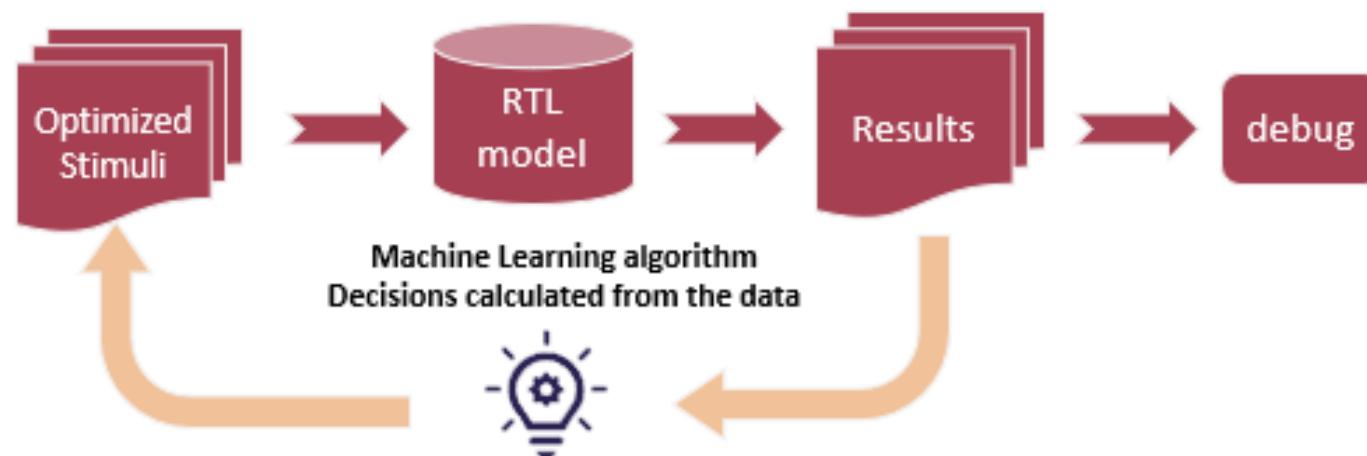


2. Optimizing Weekly regression

- Predefined high count **static** lists running on a weekly basis
 - Cover all desired scenarios
 - Maximum randomization
 - **Save compute resources**
 - Minimum runtime – enhance R&D productivity

2. Optimizing Weekly regression

- Predefined high count **static** lists running on a weekly basis
- How can AI help me?
 - Analyze static lists (“batch 0”) results → recommend new lists (“ML batches”)
 - Running ML lists based on prior runs – those lists are **dynamic**.

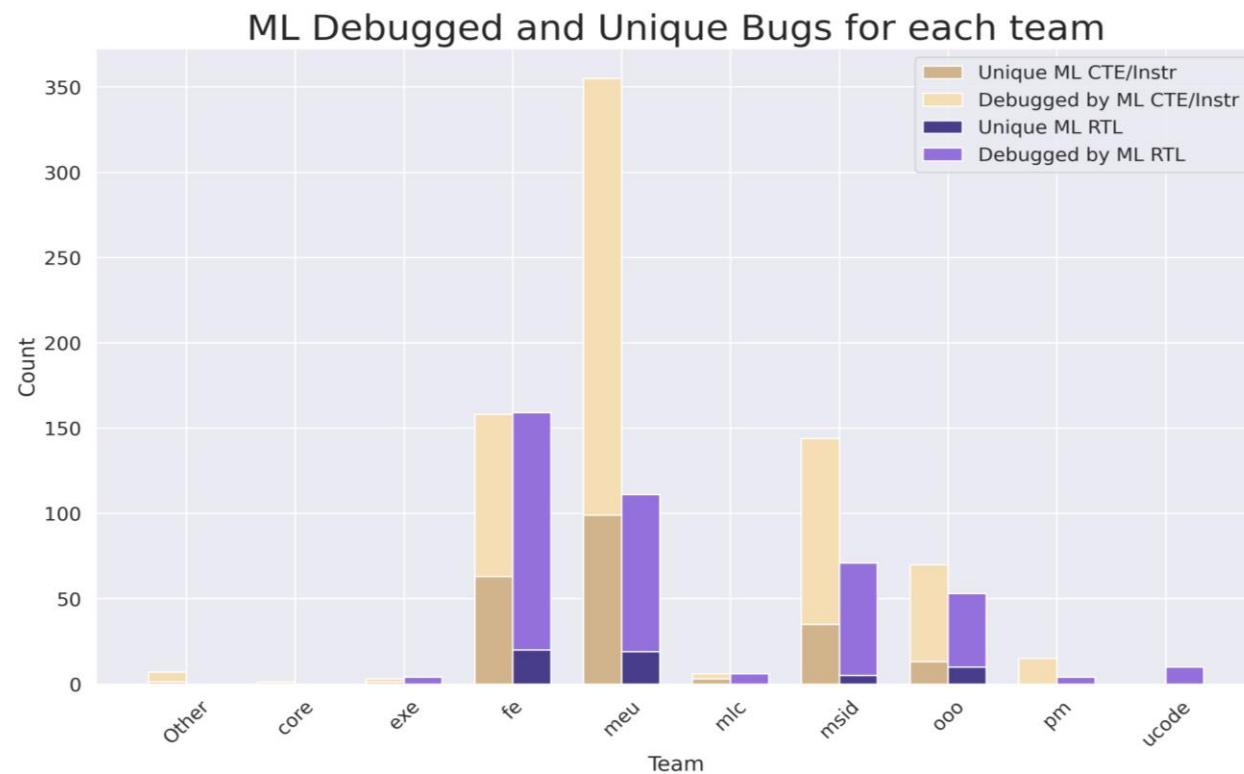


3. Intelligent stimuli generation

- Our target is to maximize DUT coverage while maintaining moderate compute resources
- We want to maximize coverage
 - Code coverage
 - Functional coverage
- We want to minimize:
 - Number of tests
 - simulation runtime

3. Intelligent stimuli generation

- How can AI help?
 - identify redundant or ineffective tests in weekly regression
 - Suggest configurations to enhance coverage (non hit/ low count scenarios)
 - Analysis of the total coverage map, point out blind spots that otherwise might be skipped



4. Reproducing rare failures

- During **Closure**  stage we focus on debug and coverage gaps
 - Debugging rare failure clusters is challenging:
 - Low test count
 - less data for ML analysis
 - Not reproduce in each weekly regression
 - lower priority for debug
 - The rarest failures sometimes “hide” a unique RTL bug

4. Reproducing rare failures

- During **Closure**  stage we focus on debug and coverage gaps
 - Using AI we can:
 - review historical regression data
 - Analyze the data and find patterns
 - Creating a list with ML configurations
 - Analyze the results to optimize ML recommendations



Prob(X)=50%

Prob(X given A=True and C=False)=75%

Input A	Input B	Input C	Input D	Target
TRUE	TRUE	FALSE	FALSE	X
TRUE	FALSE	FALSE	TRUE	X
TRUE	TRUE	FALSE	TRUE	Y
FALSE	FALSE	TRUE	FALSE	Y
FALSE	FALSE	TRUE	TRUE	Y
TRUE	TRUE	FALSE	TRUE	X

4. Reproducing rare failures

- During **Closure** 📚 stage we focus on coverage gaps
- Using AI we can:
 - review historical regression data
 - Analyze the data and find patterns
 - Creating a list with ML configurations
- Analyze the results to optimize ML recommendations



Examples from PNC:

Vanishing buckets lists hit RTL bugs in FE_000 and CTE bug in MSID

Prob(X)=50%

Prob(X given A=True and C=False)=75%

Input A	Input B	Input C	Input D	Target
TRUE	TRUE	FALSE	FALSE	X
TRUE	FALSE	FALSE	TRUE	X
TRUE	TRUE	FALSE	TRUE	Y
FALSE	FALSE	TRUE	FALSE	Y
FALSE	FALSE	TRUE	TRUE	Y
TRUE	TRUE	FALSE	TRUE	X

summary

- CPU R&D and verification is a highly complicated project
- As technology advances, the challenge of achieving high confidence in the product prior to TAPE-IN is increasing
- AI is optimizing many of the verification tasks to enhance productivity and better use our resources.

Thank you!

- annieda666@yahoo.com
- [linkedin.com/in/lotem-dallal/](https://www.linkedin.com/in/lotem-dallal/)

